

T2. Data Structures and Basic Operations

J. E. Roman

Departament de Sistemes Informàtics i Computació
Universitat Politècnica de València

Year 2021-2022



UNIVERSITAT
POLITÀCNICA
DE VALÈNCIA

1

Contents

- 1** Fields, Vectors
 - Storage Formats
 - Basic Operations
 - Orthogonalization
- 2** Sparse Matrices
 - Matrix Properties and Types
 - Storage Formats
 - Basic Operations
- 3** Orderings and Partitioning
 - Meshes, Graphs
 - Orderings
 - Partitioning

2

Section 1

Fields, Vectors

- Storage Formats
- Basic Operations
- Orthogonalization

3

Representation of Fields: Vectors

Need to represent the different magnitudes defined in the domain (fields)

The **discrete version of the field** is represented associating numeric values to mesh entities:

- Nodes (vertices) - e.g. finite differences
- Cells (elements) - e.g. finite volumes
- Sometimes also in edges or faces

Scalar fields (e.g. pressure) are represented with 1 value, whereas vector fields (e.g. velocity) are represented with 2 or 3


It is important to keep the correspondence between elements of the vector and the associated mesh entities

4

Natural Ordering

Usually the “natural ordering” is employed

- From left to right, from the bottom up



$$u = [u_{11}, u_{21}, u_{31}, u_{41}, u_{12}, \dots, u_{34}, u_{44}]^T$$

In this case, $u \in \mathbb{R}^n$ with $n = 16$


- In vector fields, $u \in \mathbb{R}^{dn}$ with $d = 2, 3$

5

Parallel Vectors

Each process locally stores a subvector


- It is convenient that the size of the subvector n_i be approximately equal in all processes $i = 0, \dots, p - 1$
- Local indices $1, \dots, n_i$ correspond to global indices



6

Parallel Vectors: Example

We want to distribute the vector among several processes:



Block distribution, $p = 2$

$$\rightarrow P_0 = \{1, 2, 3, 4, \dots, 8\}, P_1 = \{9, 10, 11, 12, \dots, 16\}$$

Block distribution, $p = 3$

$$\rightarrow P_0 = \{1, 2, \dots, 6\}, P_1 = \{7, 8, \dots, 11\}, P_2 = \{12, 13, \dots, 16\}$$

Block cyclic distribution, $n_b = 4, p = 2$

$$\rightarrow P_0 = \{1, 2, 3, 4, 9, 10, 11, 12\}, P_1 = \{5, 6, 7, 8, 13, 14, 15, 16\}$$

Block cyclic distribution, $n_b = 4, p = 3$

$$\rightarrow P_0 = \{1, 2, 3, 4, 13, 14, 15, 16\}, P_1 = \{5, 6, 7, 8\}, P_2 = \{9, 10, 11, 12\}$$

7

Vector Operations

Let $x, y, z \in \mathbb{R}^n$ and $\alpha, r \in \mathbb{R}$

Sum and product by a scalar

- Sum: $z = x + y, \quad z_i = x_i + y_i$
- Scaling: $y = \alpha x, \quad y_i = \alpha x_i$
- Combined operation (AXPY): $y = \alpha x + y$
- Trivially parallelizable

Vector **dot product** (also *scalar product* or *inner product*)

- $r = x^T y$

$$r = \sum_{i=1}^n x_i y_i$$

- In parallel: reduction of the partial sums

8

Vector Norm

A vector norm is a real function $\|\cdot\|$ satisfying

1 $\|x\| \geq 0 \quad \forall x \in \mathbb{R}^n, \quad \|x\| = 0 \Leftrightarrow x = 0$

2 $\|x + y\| \leq \|x\| + \|y\| \quad \forall x, y \in \mathbb{R}^n$

3 $\|\alpha x\| = |\alpha| \|x\| \quad \forall \alpha \in \mathbb{R} \quad \forall x \in \mathbb{R}^n$

Vector 2-norm or Euclidean norm

$$\|x\|_2 = \sqrt{x_1^2 + x_2^2 + \cdots + x_n^2} = \sqrt{x^T x}$$

Other: 1-norm, ∞ -norm

Normalization: to scale a vector so that its norm is 1

$$\tilde{x} = \frac{x}{\|x\|_2}$$

One can check that $\|\tilde{x}\|_2 = 1$

9

Multi-Vectors

Given k vectors $\{v_1, v_2, \dots, v_k\}$, they are sometimes combined in a matrix

$$V = \begin{bmatrix} | & | & & | \\ v_1 & v_2 & \cdots & v_k \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times k}$$

Operations

■ Multiple AXPY: let $s = [s_1, s_2, \dots, s_k]^T$

$$y = y + Vs = y + \sum_{i=1}^k s_i v_i$$

■ Multiple dot product: $s = V^T x$ where $s_i = v_i^T x$

10

Vector Subspaces

$\mathcal{V} = \text{span}\{v_1, v_2, \dots, v_k\}$ is the subspace spanned by the vectors v_1, \dots, v_k

- If the v_i 's are linearly independent, then they form a **basis** of \mathcal{V} and $\dim(\mathcal{V}) = k$
- Otherwise, $\dim(\mathcal{V}) < k$

Let $V = [v_1, v_2, \dots, v_k]$ and $s = [s_1, s_2, \dots, s_k]^T$

- $y = Vs$ is a linear combination of the v_i 's
- Therefore, $y \in \mathcal{V}$

Let $t = [t_1, t_2, \dots, t_k]^T$ and $s = Ht$

- $y = Vs = VHt = Wt$
- Change of basis: the columns of $W = VH$ generate the same subspace (if H is non-singular)

11

Orthogonalization of Vectors

Orthogonal vectors

- Two vectors $x, y \in \mathbb{R}^n, x \neq 0, y \neq 0$ are orthogonal if $x^T y = 0$
- A set $\{q_1, q_2, \dots, q_k\}$ is orthogonal if $q_i^T q_j = 0$ for $i \neq j$
- We say it is orthonormal if in addition $q_i^T q_i = 1$
- In matrix form: $Q^T Q = I$ (I = identity matrix)

Problem: obtain an orthonormal basis of a subspace

- Input: V such that $\mathcal{V} = \text{span}\{v_1, v_2, \dots, v_k\}$
- Output: Q such that the subspace generated by the q_i 's is also \mathcal{V} (change of basis) and they are orthonormal

12

Gram-Schmidt: Idea


The Gram-Schmidt method generates the q_i 's one by one

First vector:

- Normalize: $q_1 = v_1 / \|v_1\|_2$ (assuming $\|v_1\| \neq 0$)

Second vector:

- Subtract its **projection** on q_1 : $\hat{v}_2 = v_2 - (q_1^T v_2)q_1$
- Normalize: $q_2 = \hat{v}_2 / \|\hat{v}_2\|_2$ (assuming $\|\hat{v}_2\| \neq 0$)



In general: subtract from v_j its projection on q_1, q_2, \dots, q_{j-1}

13

Classical Gram-Schmidt

Given a basis $\{v_1, v_2, \dots, v_k\}$ obtain an orthonormal basis of the same subspace $\{q_1, q_2, \dots, q_k\}$

CGS: Classical Gram-Schmidt

```
 $r_{11} = \|v_1\|_2$  (if  $r_{11} = 0$  abort)
 $q_1 = v_1 / r_{11}$ 
for  $j = 2, \dots, k$ 
  for  $i = 1, \dots, j - 1$ 
     $r_{i,j} = q_i^T v_j$ 
  end
   $\hat{v}_j = v_j$ 
  for  $i = 1, \dots, j - 1$ 
     $\hat{v}_j = \hat{v}_j - r_{i,j} q_i$ 
  end
   $r_{j,j} = \|\hat{v}_j\|$  (if  $r_{j,j} = 0$  abort)
   $q_j = \hat{v}_j / r_{j,j}$ 
end
```

14

Modified Gram-Schmidt

MGS: Modified Gram-Schmidt

```
 $r_{11} = \|v_1\|_2$  (if  $r_{11} = 0$  abort)  
 $q_1 = v_1/r_{11}$   
for  $j = 2, \dots, k$   
   $\hat{v}_j = v_j$   
  for  $i = 1, \dots, j - 1$   
     $r_{i,j} = q_i^T \hat{v}_j$   
     $\hat{v}_j = \hat{v}_j - r_{i,j}q_i$   
  end  
   $r_{j,j} = \|\hat{v}_j\|$  (if  $r_{j,j} = 0$  abort)  
   $q_j = \hat{v}_j/r_{j,j}$   
end
```

MGS numerically more stable than CGS

CGS more efficient in parallel: the scalar products

$r_{1:j-1,j} = Q_{j-1}^T v_j$ are merged in a single communication

15

QR Decomposition

The result of Gram-Schmidt satisfies $v_j = \sum_{i=1}^j r_{i,j}q_i$

In matrix notation: $V = QR$, being R upper triangular

$$\left[\begin{array}{c|c|c|c} v_1 & v_2 & \cdots & v_k \end{array} \right] = \left[\begin{array}{c|c|c|c} q_1 & q_2 & \cdots & q_k \end{array} \right] \left[\begin{array}{cccc} r_{11} & r_{12} & \cdots & r_{1,k} \\ & r_{22} & \cdots & r_{2,k} \\ & & \ddots & \vdots \\ & & & r_{k,k} \end{array} \right]$$

R is the matrix of the change of basis

16

Section 2

Sparse Matrices

- Matrix Properties and Types
- Storage Formats
- Basic Operations

17

Symmetric Matrices

A square matrix $A \in \mathbb{R}^{n \times n}$ is symmetric if $a_{ij} = a_{ji}$

$$A = A^T$$

Symmetric positive-definite matrix: satisfies

$$x^T A x > 0 \quad \forall x \in \mathbb{R}^n, x \neq 0$$

Characterization:

- a matrix is SPD iff $\det(A_{1:i,1:i}) > 0, \quad 1 \leq i \leq n$

Example:

$$A = \begin{bmatrix} 5 & 2 & -1 \\ 2 & 2 & 1 \\ -1 & 1 & 3 \end{bmatrix} \quad \begin{aligned} \det \left(\begin{bmatrix} 5 \end{bmatrix} \right) &= 5 > 0 \\ \det \left(\begin{bmatrix} 5 & 2 \\ 2 & 2 \end{bmatrix} \right) &= 6 > 0 \\ \det(A) &= 7 > 0 \end{aligned}$$

18

Orthogonal Matrices

A square matrix $Q \in \mathbb{R}^{n \times n}$ is **orthogonal** if its columns are orthonormal

$$Q^T Q = I$$

Example: $Q = \begin{bmatrix} 1/2 & 0 & \sqrt{3}/2 \\ 0 & 1 & 0 \\ -\sqrt{3}/2 & 0 & 1/2 \end{bmatrix}$

Properties:

- $Q^{-1} = Q^T$
- $\|Qx\|_2 = \|x\|_2$

Permutation matrix: orthogonal matrix whose elements are zero except a 1 in each row and column

Example: $P = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$

Sometimes represented as an index vector
 $p = [2, 3, 1, 4]$

19

Diagonally Dominant Matrices

We say a matrix is **diagonally dominant**

by rows:

$$|a_{ii}| \geq \sum_{j \neq i} |a_{ij}|, \quad i = 1, \dots, n$$

by columns:

$$|a_{jj}| \geq \sum_{i \neq j} |a_{ij}|, \quad j = 1, \dots, n$$

In case of strict inequality, we say **strictly DD**

$$|a_{ii}| > \sum_{j \neq i} |a_{ij}|, \quad i = 1, \dots, n$$

Examples:

$$A = \begin{bmatrix} -3 & 0 & 0 & -1 \\ 2 & 6 & 3 & 0 \\ 2 & 1 & -6 & -2 \\ 3 & 0 & 0 & -4 \end{bmatrix}$$


Strictly DD by rows

$$T = \begin{bmatrix} 2 & -1 & & \\ -1 & 2 & -1 & \\ & -1 & 2 & -1 \\ & & -1 & 2 \end{bmatrix}$$


Diagonally dominant

20


Triangular Matrices



Upper triangular
 $a_{ij} = 0$ for $i > j$



Lower triangular
 $a_{ij} = 0$ for $i < j$




Upper Hessenberg
 $a_{ij} = 0$ for $i > j + 1$

Band matrix

$a_{ij} = 0$ if $|i - j| > \beta$ (β = bandwidth)



Particular cases: diagonal ($\beta=0$), tridiagonal ($\beta=1$), pentadiagonal ($\beta=2$)



21

Sparse Matrices

In mesh-based applications, the percentage of nonzero elements is usually very small ($<1\%$)



The **structure** or **nonzero pattern** depends on

- The connectivity between unknowns (according to discretization)
- The ordering that has been chosen for the unknowns

22

Sparse Storage Formats

Special data structures, to:

- Reduce the memory requirements
- Reduce the cost of operations

There are many different formats – intended to optimize the **most frequent operations**

- Matrix-vector product
The most frequent one, requires maximum efficiency
- Extraction of the diagonal
- Factorization
Produces “fill-in”, new nonzero elements appear
- Other operations: addition, matrix-matrix product
- Insert/delete elements
Infrequent, constant nonzero pattern

23

Coordinate Format (COO)

Uses three arrays:

A	nz	numeric values
I	nz	row indices
J	nz	column indices

(nz is the number of nonzero elements)

A	a_{11}	a_{13}	a_{21}	a_{22}	a_{24}	a_{33}	a_{44}	a_{45}	a_{52}	a_{53}	a_{55}
I	1	1	2	2	2	3	4	4	5	5	5
J	1	3	1	2	4	3	4	5	2	3	5

$$\begin{bmatrix} a_{11} & & a_{13} & & & & & & & & & \\ & a_{21} & a_{22} & & a_{24} & & & & & & & \\ & & & a_{33} & & & & & & & & \\ & & & & a_{44} & a_{45} & & & & & & \\ & & & & & & a_{52} & a_{53} & & a_{55} & & \end{bmatrix}$$

- Very efficient sequential access
- Inefficient element search


24

Compressed Sparse Row (CSR)

Uses three arrays:

A	nz	numeric values
I	n+1	pointers to row start
J	nz	column indices

(n is the number of rows of the matrix)




- Row indices are not stored explicitly
- $I[k]$ indicates the position in which row k starts

25


Variants

Compressed Sparse Column (CSC)

- Analog of CSR, but sorted by columns
- $J[k]$ indicates the position in which column k starts



Modified Sparse Row (MSR)



The main diagonal is stored separately

26

Sparse Matrix Operations

Let $A, B \in \mathbb{R}^{m \times n}$, $v \in \mathbb{R}^n$, $w \in \mathbb{R}^m$ and $\alpha \in \mathbb{R}$

Sum and product by a scalar

- Matrix AXPY: $B = \alpha A + B$
- Difficult if the pattern of A is not a subset of B 's

Product of matrices, $A = CD$

- The dimensions of C and D must be consistent
- Very uncommon in sparse matrices

Matrix-vector product, $w = Av$

- Fundamental operation in iterative methods

27

Matrix Norm

Definition is analog of vector norm

Frobenius norm: $\|A\|_F = \sqrt{\sum_{i,j} a_{i,j}^2}$

Matrix 2-Norm (also 1-norm, ∞ -norm)

- Induced from the vector 2-norm

$$\|A\|_2 = \max_{\|y\|_2=1} \|Ay\|_2$$

- Intuitively: maximum produced elongation
- Consistent norms: $\|Ax\|_2 \leq \|A\|_2 \|x\|_2$

Condition number

$$\kappa_2(A) = \|A\|_2 \|A^{-1}\|_2, \quad \kappa_2(A) \geq 1$$

28

Matrix-Vector Product Algorithm

$$w = Av$$

Matrix-vector product by rows

```
 $w = [0, 0, \dots, 0]^T$   
for  $i = 1, \dots, m$   
  for  $j = 1, \dots, n$   
     $w[i] += A[i][j] * v[j]$   
  end  
end
```

Matrix-vector prod. COO

```
 $w = [0, 0, \dots, 0]^T$   
for  $k = 1, \dots, nz$   
   $w[I[k]] += A[k] * v[J[k]]$   
end
```

Matrix-vector prod. CSR

```
 $w = [0, 0, \dots, 0]^T$   
for  $i = 1, \dots, m$   
  for  $k = I[i], \dots, I[i+1] - 1$   
     $w[i] += A[k] * v[J[k]]$   
  end  
end
```

29

Section 3

Orderings and Partitioning



- Meshes, Graphs
- Orderings
- Partitioning

30

Sparse Pattern

The nonzero element pattern depends on

- The connectivity in the mesh
- The chosen ordering



The ordering can be chosen to improve some structural property of the matrix (bandwidth, fill-in, ...)


31

Adjacency Graph


We use graph theory to study the sparsity structure

Adjacency Graph of a sparse matrix $A \in \mathbb{R}^{n \times n}$, $G(A) = (V, E)$

- The n vertices of V represent the unknowns
- The edges E represent binary relations among them
 - An edge $(v_i, v_j) \in E$ if $a_{ij} \neq 0$
 - Represents the presence of unknown j in equation i



Non-symmetric; Directed graph




Symmetric; Undirected graph


32

Correspondence between Mesh and Graph

Discretization mesh



Choose an ordering



Adjacency graph

System matrix

$$A = \begin{bmatrix} 4 & -1 & & & & \\ -1 & 4 & -1 & & & \\ & -1 & 4 & & -1 & \\ -1 & & & 4 & -1 & -1 \\ & -1 & & -1 & 4 & -1 \\ & & -1 & & -1 & 4 \\ & & & -1 & & 4 & -1 \\ & & & & -1 & -1 & 4 \end{bmatrix}$$

33

Symmetric Permutation

Given a system $Ax = b$, and let P be a permutation matrix

- Row permutation $PAx = Pb$, reorders equations
- Column permutation $AP^T Px = b$, exchanges unknowns
- Symmetric permutation $(PAP^T)Px = Pb$, modifies both

Example with permutation $p = [3, 2, 4, 1]$:

$$A = \begin{bmatrix} a_{11} & a_{12} & & \\ a_{21} & a_{22} & a_{23} & \\ & a_{32} & a_{33} & a_{34} \\ & & a_{43} & a_{44} \end{bmatrix}, \quad \begin{bmatrix} a_{33} & a_{32} & a_{34} & \\ a_{23} & a_{22} & & a_{21} \\ a_{43} & & a_{44} & \\ & a_{12} & & a_{11} \end{bmatrix} \begin{bmatrix} x_3 \\ x_2 \\ x_4 \\ x_1 \end{bmatrix} = \begin{bmatrix} b_3 \\ b_2 \\ b_4 \\ b_1 \end{bmatrix}$$



Adjacency graph




Graph after permutation

→ the graph vertices are relabelled without modifying edges

34


Red-Black Ordering

Natural Ordering




From left to right and
from the bottom up

Example, pattern for
 $n = 6 \rightarrow$




Red-Black Ordering



1. Color the nodes in a way
that neighbors have different
color


2. Number the red nodes
first, then the black ones
 $p = [1, 3, 5, 7, 9, 2, 4, 6, 8]$




35

Other Orderings

The main purpose of reordering is to improve the structural
properties of the matrix



Reverse Cuthill-McKee,
minimizes the bandwidth



Minimum degree, reduces fill-in
in the factorizations


Other:

- *Nested dissection*, also reduces fill-in
- Partitioning: minimize communications in parallel

36

Parallel Matrix-Vector Product

Block-row distribution with natural ordering




The volume of communication in parallel depends on:

- Nonzero elements outside the diagonal block
- Number of edge cuts (go from one subdomain to another)

37

Mesh Partitioning

Also in unstructured meshes




Each process sends the unknowns lying on the interface to the corresponding neighboring processes, and receives in turn from them

38

Graph Partitioning

Goal: divide the mesh in regions of **similar size** to be assigned to each process, trying to reduce the **communication cost**



Types of techniques:

- Iterative exchange
- Recursive Bisection
- Multilevel


Most algorithms are based on **adjacency graphs**: undirected graph, with or without weights, $G = (N, E, W_N, W_E)$

- Balance the sum of weights W_N in each partition
- Minimize the sum of weights W_E from one partition to another

39

Parallel Vectors

- For best efficiency, the position of elements to send is precomputed
- It is often necessary to have a local representation of the vector including the received values (*ghost values*)



40